

# Monte Carlo Pricing Optimization (CPU vs GPU)

Parallel Programming Final Project — Group 23

Members: David Lu (T14902116), 蔡琦皇 (P13922006), 楊翊廷 (R14944018), 詹湧翔 (B13201026)

## Introduction

Monte Carlo (MC) simulation is one of the most flexible approaches for option pricing, especially for products without closed-form solutions or with complex payoffs (e.g., Asian options) and higher-dimensional problems (e.g., basket/multi-asset options). However, MC is computationally expensive: achieving low statistical error typically requires **millions to billions of paths** and (for path-dependent products) **hundreds of time steps**, which can be too slow on CPUs for latency-sensitive use cases.

This project implements MC option pricing with **three levels of parallelism**:

- **CPU OpenMP** (multi-core parallel paths)
- **Single-GPU CUDA** (massively parallel path simulation)
- **Multi-GPU CUDA** (distribute paths across multiple GPUs for basket pricing)

In addition, we calibrate simple GBM parameters from real historical data ([data/](#)) and benchmark performance/scalability under different workloads.

## Method

### 1) Stochastic model (GBM)

We simulate the underlying asset price using Geometric Brownian Motion (GBM):

$$S_{t+\Delta t} = S_t \cdot \exp\left(\left(r - \frac{1}{2}\sigma^2\right)\Delta t + \sigma\sqrt{\Delta t}, Z\right), \quad Z \sim \mathcal{N}(0, 1)$$

For multi-asset baskets, we generate **correlated Gaussian vectors**  $Z$  using an equicorrelation covariance structure and **Cholesky decomposition**  $\sigma = LL^\top$ , then apply  $y = Lz$ .

### 2) Payoffs implemented

This repo implements three core payoffs (plus different experiment groupings):

- **European Call (single asset)**: payoff depends only on terminal price  $S_T$ .
  - Validation: compare MC estimate vs **Black-Scholes closed-form**.
- **Asian Arithmetic Call (single asset)**: payoff depends on arithmetic average along the path.
- **Basket European Call (multi-asset)**: payoff depends on arithmetic mean of terminal prices across assets; correlation handled via Cholesky.

CPU implementation: [src/cpu/mc\\_pricer.cpp](#)

GPU implementation: [src/gpu/gpu\\_pricer.cu](#) + core kernel/workspace in [src/gpu/mc\\_pricer.cu](#)

### 3) Parallelization strategy

- **CPU (OpenMP)**: each thread simulates independent paths (per-thread RNG seeded by thread id).

- Implementation: OpenMP parallel loop in `src/cpu/mc_pricer.cpp`.
- **GPU (CUDA):**
  - RNG: per-thread cuRAND Philox states (`curandStatePhilox4_32_10_t`).
  - Kernel design: grid-stride loop where each CUDA thread processes many paths.
  - Accumulation: atomic accumulation of sum and sum of squares to compute mean and standard error.
  - Workspace reuse: allocate device buffers once (states/sum/sum2) and reuse across calls.
  - Implementation: `src/gpu/mc_pricer.cu` (kernel `mc_kernel_general`, workspace `GpuWorkspace`).
- **Multi-GPU (basket):**
  - Split total paths across devices; each GPU runs the basket kernel on its portion.
  - Combine partial means/variances to produce final mean and standard error.
  - Implementation: `src/gpu/gpu_pricer.cu` (multi-GPU path splitting + merge).

#### 4) Parameter calibration from real data

We calibrate basic GBM parameters from historical CSVs:

- $S_0$ : last close price
- $\sigma$ : annualized volatility from daily log returns
- $\rho$ : average off-diagonal correlation for the multi-asset set

Implementation: `src/tools/calibrate_from_data.py`

Outputs a shell script `params.sh` used by `src/cpu/run_all.sh` and `src/gpu/run_all.sh`.

#### Experiment

All experiments are scripted and write results as CSV rows with timing and standard error, then generate plots.

#### 1) How experiments are run

- **GPU experiments:** `src/gpu/run_experiments.sh`
- **CPU experiments:** `src/cpu/run_experiments.sh`
- **Run everything + plot:** `src/run_all_experiments.sh`
- **Plotting:** `src/plot_experiments.py` (reads `src/experiments/results/*.csv`, writes `src/experiments/plots/*.png`)

#### 2) Experiment dimensions (what we vary)

From the proposal and the actual scripts, the repo benchmarks:

- **Paths scaling:** vary number of Monte Carlo paths  $N$  (GPU goes up to  $10^9$  paths in script; CPU uses smaller  $N$ ).
- **Steps scaling:** vary time steps (e.g., 64, 128, 252, 512) for Asian/basket discretization.
- **CPU threads scaling** (OpenMP): vary thread count for basket.
- **Multi-GPU scaling** (basket): vary number of GPUs (1/2/4) and measure speedup.

- **Assets scaling** (basket): vary number of assets (2/4/8/16/32) to show dimensionality cost.
- **GPU configuration tuning**: block size and blocks-per-SM (occupancy) tuning.
- **Option type comparison**: compare European vs Asian vs Basket under same path counts.
- **Direct GPU vs CPU comparison**: for multiple option types under aligned path counts.

### 3) Outputs

Results are written to:

- `src/experiments/results/*.csv`

Plots are written to:

- `src/experiments/plots/*.png`

Each CSV row includes (examples): engine, type, workers (threads or GPUs), paths, steps, assets, rho, S0/K/r/sigma/T, price, std\_error, time\_ms (and GPU tuning fields for block size).

### Conclusion

Based on the project slides and the implemented experiments:

- **GPU acceleration is substantial** for MC pricing because paths are embarrassingly parallel.
- **Complexity increases significantly** as we move from single-asset to multi-asset basket pricing due to correlation handling (Cholesky and matrix-vector operations) and higher dimensional simulation.
- **Speedup characteristics differ by product:**
  - Asian option shows a relatively consistent GPU speedup in the reported results.
  - European option speedup depends more on path count (GPU becomes more favorable for larger workloads).
- **Best GPU configuration (from block tuning in slides):**
  - Block size  $\approx 256$
  - Occupancy setting  $\approx 2 \text{ blocks per SM}$
- Overall, the repo demonstrates that CUDA (and multi-GPU for basket) can reduce runtime dramatically while keeping pricing accuracy consistent (standard error decreases with more paths, and European can be validated against Black–Scholes).

### Work Distribution

- **David Lu (T14902116)**: Presentation
- **蔡琦皇 (P13922006)**: CPU baseline implementation (OpenMP Monte Carlo pricer), single-asset implementation, CPU experiment scripts.
- **楊翊廷 (R14944018)**: Experiment design/automation and results visualization (CSV schema, plotting scripts, figure generation), multi-asset implementation, Data calibration tooling from real market CSVs and integration into run pipelines
- **詹清翔 (B13201026)**: report generation.

ref.

- 1.An introduction to Monte Carlo methods [<https://arxiv.org/pdf/1404.0209.pdf>]
- 2.GitHub [<https://github.com/Yiting1022/pp-final.git>]